

# Mnist

June 26, 2025

## 1 Ahmad Siddiqui

## 2 Manual task 1 cateory 2

### 2.1 Question

1. OvO vs. OvR:

One-vs-One (OvO) and One-vs-Rest (OvR) are strategies for extending binary classifiers to multi-class problems.

Feature	One-vs-Rest (OvR)	One-vs-One (OvO)
<b>Number of classifiers</b>	$n$	$n(n - 1)/2$
<b>Each classifier trained on</b>	All data, with one class vs rest	Only data from two classes
<b>Training time</b>	Faster (fewer classifiers)	Slower (more classifiers)
<b>Prediction</b>	Highest decision score among classifiers	Majority vote among all pairwise classifiers
<b>Memory usage</b>	Lower	Higher
<b>Good for many classes</b>	Yes	No (quadratic growth)
<b>Boundary granularity</b>	Coarser	Finer
<b>Score interpretability</b>	Easier (direct per-class scores)	Harder (votes must be aggregated)
<b>Supported in sklearn</b>	LogisticRegression, RidgeClassifier, etc.	SVC (default with kernel), GaussianNB, etc.
<b>Typical use case</b>	Linear models, high-class problems	SVMs, low to moderate number of classes

2. SGD vs RF classifier

Feature / Aspect	SGDClassifier	RandomForestClassifier
<b>Type</b>	Linear model trained with stochastic gradient descent	Ensemble of decision trees (bagging)
<b>Model Complexity</b>	Linear (or linear + kernel trick manually)	Non-linear, high-capacity

Feature / Aspect	SGDClassifier	RandomForestClassifier
<b>Training Speed</b>	<b>Very fast</b> , scalable to millions of samples	Slower, especially with many trees or features
<b>Prediction Speed</b>	<b>Very fast</b> (single dot product)	Slower (tree traversal per tree, many trees)
<b>Memory Usage</b>	Low	High (stores many full trees)
<b>Handles Multiclass</b>	Yes (ovr, multinomial)	Yes (natively)
<b>Works with Sparse Data</b>	<b>Yes (natively)</b>	No (must densify input)
<b>Regularization</b>	Supports L1, L2, ElasticNet	No direct regularization, prone to overfitting
<b>Handles Non-linear Data</b>	Poorly unless kernelized or feature engineered	<b>Very well</b>
<b>Overfitting Risk</b>	Low (but high bias)	Medium to high (trees can memorize noise)
<b>Interpretability</b>	High (weights per feature)	Medium-low (can use feature importances)
<b>Feature Scaling Required</b>	<b>Yes</b> (mandatory for convergence)	<b>No</b>
<b>Parallelization</b>	Difficult (single-pass)	<b>Yes</b> (trees can be grown in parallel)
<b>Handles Missing Values</b>	No (must impute manually)	Yes (can tolerate missing splits)
<b>Typical Use Cases</b>	Text classification, online learning, large-scale ML	Tabular data, feature-rich problems, small-medium datasets

```
[1]: import sys
assert sys.version_info >= (3, 5)
import sklearn
import numpy as np
import os
np.random.seed(42)

%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt
mpl.rc('axes', labelsizes=14)
mpl.rc('xtick', labelsizes=12)
mpl.rc('ytick', labelsizes=12)

# Where to save the figures
PROJECT_ROOT_DIR = "./content"
CHAPTER_ID = "assignment"
IMAGES_PATH = os.path.join(PROJECT_ROOT_DIR, "images", CHAPTER_ID)
os.makedirs(IMAGES_PATH, exist_ok=True)
```

```
def save_fig(fig_id, tight_layout=True, fig_extension="png", resolution=300):
    path = os.path.join(IMAGES_PATH, fig_id + "." + fig_extension)
    print("Saving figure", fig_id)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)
```

```
[2]: from sklearn.datasets import fetch_openml
mnist = fetch_openml('mnist_784', version=1, as_frame=False)
mnist.keys()
```

```
[2]: dict_keys(['data', 'target', 'frame', 'categories', 'feature_names',
'target_names', 'DESCR', 'details', 'url'])
```

```
[3]: X, y = mnist["data"], mnist["target"]
X.shape
```

```
[3]: (70000, 784)
```

```
[4]: y.shape
```

```
[4]: (70000,)
```

```
[5]: 28 * 28
```

```
[5]: 784
```

```
[6]: %matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt

some_digit = X[0]
some_digit_image = some_digit.reshape(28, 28)
plt.imshow(some_digit_image, cmap=mpl.cm.binary)
plt.axis("off")

save_fig("some_digit_plot")
plt.show()
```

Saving figure some\_digit\_plot



```
[7]: y[0]
```

```
[7]: '5'
```

```
[8]: y = y.astype(np.uint8)
```

```
[9]: def plot_digit(data):  
    image = data.reshape(28, 28)  
    plt.imshow(image, cmap = mpl.cm.binary,  
               interpolation="nearest")  
    plt.axis("off")
```

```
[ ]: y[0]
```

```
[ ]: 5
```

```
[12]: X_train, X_test, y_train, y_test = X[:60000], X[60000:], y[:60000], y[60000:]
```

### 3 Training a Binary Classifier

```
[13]: y_train_5 = (y_train == 5)
      y_test_5 = (y_test == 5)
```

```
[14]: from sklearn.linear_model import SGDClassifier

      sgd_clf = SGDClassifier(max_iter=1000, tol=1e-3, random_state=42)
      sgd_clf.fit(X_train, y_train_5)
```

```
[14]: SGDClassifier(random_state=42)
```

```
[15]: sgd_clf.predict([some_digit])
```

```
[15]: array([ True])
```

```
[16]: from sklearn.model_selection import cross_val_score
      cross_val_score(sgd_clf, X_train, y_train_5, cv=3, scoring="accuracy")
```

```
[16]: array([0.95035, 0.96035, 0.9604 ])
```

#### 3.1 Confusion Matrix

```
[17]: from sklearn.model_selection import cross_val_predict

      y_train_pred = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3)
```

```
[18]: from sklearn.metrics import confusion_matrix

      confusion_matrix(y_train_5, y_train_pred)
```


```
[18]: array([[53892,   687],
          [ 1891,  3530]])
```

```
[26]: # To plot the confusion matrix nicely, we can use `matshow()` from `matplotlib`
      ↪ pyplot`
      # and add labels and colorbar.

      from sklearn.metrics import ConfusionMatrixDisplay

      def plot_confusion_matrix(y_true, y_pred, classes):
          cm = confusion_matrix(y_true, y_pred)
          disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=classes)
          disp.plot(cmap=plt.cm.Blues)
          plt.title("Confusion Matrix")
          plt.show()
```

```
# Example usage (assuming you have y_train_5 and y_train_pred defined):
# You would need to determine the classes based on your binary classification
# (e.g., [False, True])
classes = [0, 1] # Assuming the classes are 0 (not 5) and 1 (is 5)
plot_confusion_matrix(y_train_5, y_train_pred, classes)
```



```
[66]: from sklearn.metrics import classification_report
print("*** Classification Report ***"*3)
print(classification_report(y_train_5, y_train_pred))
```

```
*** Classification Report ***** Classification Report ***** Classification
Report ***
```

	precision	recall	f1-score	support
False	0.97	0.99	0.98	54579
True	0.84	0.65	0.73	5421
accuracy			0.96	60000
macro avg	0.90	0.82	0.85	60000

weighted avg	0.95	0.96	0.95	60000
--------------	------	------	------	-------

### 3.2 Precision/Recall Trade-off

```
[31]: y_scores = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3,
                                     method="decision_function")

[32]: from sklearn.metrics import precision_recall_curve


precisions, recalls, thresholds = precision_recall_curve(y_train_5, y_scores)

[33]: def plot_precision_recall_vs_threshold(precisions, recalls, thresholds):
    plt.plot(thresholds, precisions[:-1], "b--", label="Precision", linewidth=2)
    plt.plot(thresholds, recalls[:-1], "g-", label="Recall", linewidth=2)
    plt.legend(loc="center right", fontsize=16) # Not shown in the book
    plt.xlabel("Threshold", fontsize=16)       # Not shown
    plt.grid(True)                             # Not shown
    plt.axis([-50000, 50000, 0, 1])           # Not shown

recall_90_precision = recalls[np.argmax(precisions >= 0.90)]
threshold_90_precision = thresholds[np.argmax(precisions >= 0.90)]

plt.figure(figsize=(8, 4))
    ↪ # Not shown
plot_precision_recall_vs_threshold(precisions, recalls, thresholds)
plt.plot([threshold_90_precision, threshold_90_precision], [0., 0.9], "r:")
    ↪ # Not shown
plt.plot([-50000, threshold_90_precision], [0.9, 0.9], "r:")
    ↪ # Not shown
plt.plot([-50000, threshold_90_precision], [recall_90_precision,
    ↪ recall_90_precision], "r:") # Not shown
plt.plot([threshold_90_precision], [0.9], "ro")
    ↪ # Not shown
plt.plot([threshold_90_precision], [recall_90_precision], "ro")
    ↪ # Not shown
save_fig("precision_recall_vs_threshold_plot")
    ↪ # Not shown
plt.show()
```

Saving figure precision\_recall\_vs\_threshold\_plot




```
[34]: def plot_precision_vs_recall(precisions, recalls):
    plt.plot(recalls, precisions, "b-", linewidth=2)
    plt.xlabel("Recall", fontsize=16)
    plt.ylabel("Precision", fontsize=16)
    plt.axis([0, 1, 0, 1])
    plt.grid(True)

    plt.figure(figsize=(8, 6))
    plot_precision_vs_recall(precisions, recalls)
    plt.plot([recall_90_precision, recall_90_precision], [0., 0.9], "r:")
    plt.plot([0.0, recall_90_precision], [0.9, 0.9], "r:")
    plt.plot([recall_90_precision], [0.9], "ro")
    save_fig("precision_vs_recall_plot")
    plt.show()
```

Saving figure precision\_vs\_recall\_plot





### 3.3 The ROC Curve

```
[40]: from sklearn.metrics import roc_curve
```


```
fpr, tpr, thresholds = roc_curve(y_train_5, y_scores)
```

```
[41]: def plot_roc_curve(fpr, tpr, label=None):
    plt.plot(fpr, tpr, linewidth=2, label=label)
    plt.plot([0, 1], [0, 1], 'k--') # dashed diagonal
    plt.axis([0, 1, 0, 1]) # Not shown in ↵
    ↪ the book
    plt.xlabel('False Positive Rate (Fall-Out)', fontsize=16) # Not shown
    plt.ylabel('True Positive Rate (Recall)', fontsize=16) # Not shown
    plt.grid(True) # Not shown

    plt.figure(figsize=(8, 6)) # Not shown
    plot_roc_curve(fpr, tpr)
    fpr_90 = fpr[np.argmax(tpr >= recall_90_precision)] # Not shown
    plt.plot([fpr_90, fpr_90], [0., recall_90_precision], "r:") # Not shown
```

```
plt.plot([0.0, fpr_90], [recall_90_precision, recall_90_precision], "r:") # ┐
    ↪ Not shown
plt.plot([fpr_90], [recall_90_precision], "ro") # Not shown
save_fig("roc_curve_plot") # Not shown
plt.show()
```

Saving figure roc\_curve\_plot



```
[42]: from sklearn.metrics import roc_auc_score

roc_auc_score(y_train_5, y_scores)
```

```
[42]: np.float64(0.9604938554008616)
```


```
[43]: from sklearn.ensemble import RandomForestClassifier
forest_clf = RandomForestClassifier(n_estimators=100, random_state=42)
y_probas_forest = cross_val_predict(forest_clf, X_train, y_train_5, cv=3,
                                   method="predict_proba")
```

```
[44]: y_scores_forest = y_probas_forest[:, 1] # score = proba of positive class
fpr_forest, tpr_forest, thresholds_forest = roc_curve(y_train_5, y_scores_forest)
```

```
[45]: recall_for_forest = tpr_forest[np.argmax(fpr_forest >= fpr_90)]

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, "b:", linewidth=2, label="SGD")
plot_roc_curve(fpr_forest, tpr_forest, "Random Forest")
plt.plot([fpr_90, fpr_90], [0., recall_90_precision], "r:")
plt.plot([0.0, fpr_90], [recall_90_precision, recall_90_precision], "r:")
plt.plot([fpr_90], [recall_90_precision], "ro")
plt.plot([fpr_90, fpr_90], [0., recall_for_forest], "r:")
plt.plot([fpr_90], [recall_for_forest], "ro")
plt.grid(True)
plt.legend(loc="lower right", fontsize=16)
save_fig("roc_curve_comparison_plot")
plt.show()
```

Saving figure roc\_curve\_comparison\_plot



```
[46]: roc_auc_score(y_train_5, y_scores_forest)
```

```
[46]: np.float64(0.9983436731328145)
```

```
[47]: y_train_pred_forest = cross_val_predict(forest_clf, X_train, y_train_5, cv=3)
precision_score(y_train_5, y_train_pred_forest)
```

```
[47]: 0.9905083315756169
```

```
[48]: recall_score(y_train_5, y_train_pred_forest)
```

```
[48]: 0.8662608374838591
```

## 4 Multilabel Classification

```
[49]: from sklearn.neighbors import KNeighborsClassifier

y_train_large = (y_train >= 7)
y_train_odd = (y_train % 2 == 1)
y_multilabel = np.c_[y_train_large, y_train_odd]

knn_clf = KNeighborsClassifier()
knn_clf.fit(X_train, y_multilabel)
```

```
[49]: KNeighborsClassifier()
```

```
[50]: knn_clf.predict([some_digit])
```

```
[50]: array([[False,  True]])
```

**Warning:** the following cell may take a very long time (possibly hours depending on your hardware).

```
[51]: y_train_knn_pred = cross_val_predict(knn_clf, X_train, y_multilabel, cv=3)
f1_score(y_multilabel, y_train_knn_pred, average="macro")
```

```
[51]: 0.9764102655606048
```


## 5 Multioutput Classification

```
[52]: noise = np.random.randint(0, 100, (len(X_train), 784))
X_train_mod = X_train + noise
noise = np.random.randint(0, 100, (len(X_test), 784))
X_test_mod = X_test + noise
y_train_mod = X_train
y_test_mod = X_test
```

```
[53]: some_index = 0
plt.subplot(121); plot_digit(X_test_mod[some_index])
plt.subplot(122); plot_digit(y_test_mod[some_index])
```


```
save_fig("noisy_digit_example_plot")  
plt.show()
```

Saving figure noisy\_digit\_example\_plot



```
[54]: knn_clf.fit(X_train_mod, y_train_mod)  
clean_digit = knn_clf.predict([X_test_mod[some_index]])  
plot_digit(clean_digit)  
save_fig("cleaned_digit_example_plot")
```

Saving figure cleaned\_digit\_example\_plot



```
[55]: from sklearn.neighbors import KNeighborsClassifier
knn_clf = KNeighborsClassifier(weights='distance', n_neighbors=4)
knn_clf.fit(X_train, y_train)
```

```
[55]: KNeighborsClassifier(n_neighbors=4, weights='distance')
```

```
[56]: y_knn_pred = knn_clf.predict(X_test)
```

```
[57]: from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_knn_pred)
```


```
[57]: 0.9714
```

```
[58]: from scipy.ndimage.interpolation import shift
def shift_digit(digit_array, dx, dy, new=0):
    return shift(digit_array.reshape(28, 28), [dy, dx], cval=new).reshape(784)

plot_digit(shift_digit(some_digit, 5, 1, new=100))
```

```
/tmp/ipython-input-58-3518175631.py:1: DeprecationWarning: Please import `shift`  
from the `scipy.ndimage` namespace; the `scipy.ndimage.interpolation` namespace  
is deprecated and will be removed in SciPy 2.0.0.
```

```
from scipy.ndimage.interpolation import shift
```



```
[59]: X_train_expanded = [X_train]
y_train_expanded = [y_train]
for dx, dy in ((1, 0), (-1, 0), (0, 1), (0, -1)):
    shifted_images = np.apply_along_axis(shift_digit, axis=1, arr=X_train,
    ↪dx=dx, dy=dy)
    X_train_expanded.append(shifted_images)
    y_train_expanded.append(y_train)

X_train_expanded = np.concatenate(X_train_expanded)
y_train_expanded = np.concatenate(y_train_expanded)
X_train_expanded.shape, y_train_expanded.shape
```

```
[59]: ((300000, 784), (300000,))
```

```
[60]: knn_clf.fit(X_train_expanded, y_train_expanded)
```

```
[60]: KNeighborsClassifier(n_neighbors=4, weights='distance')
```

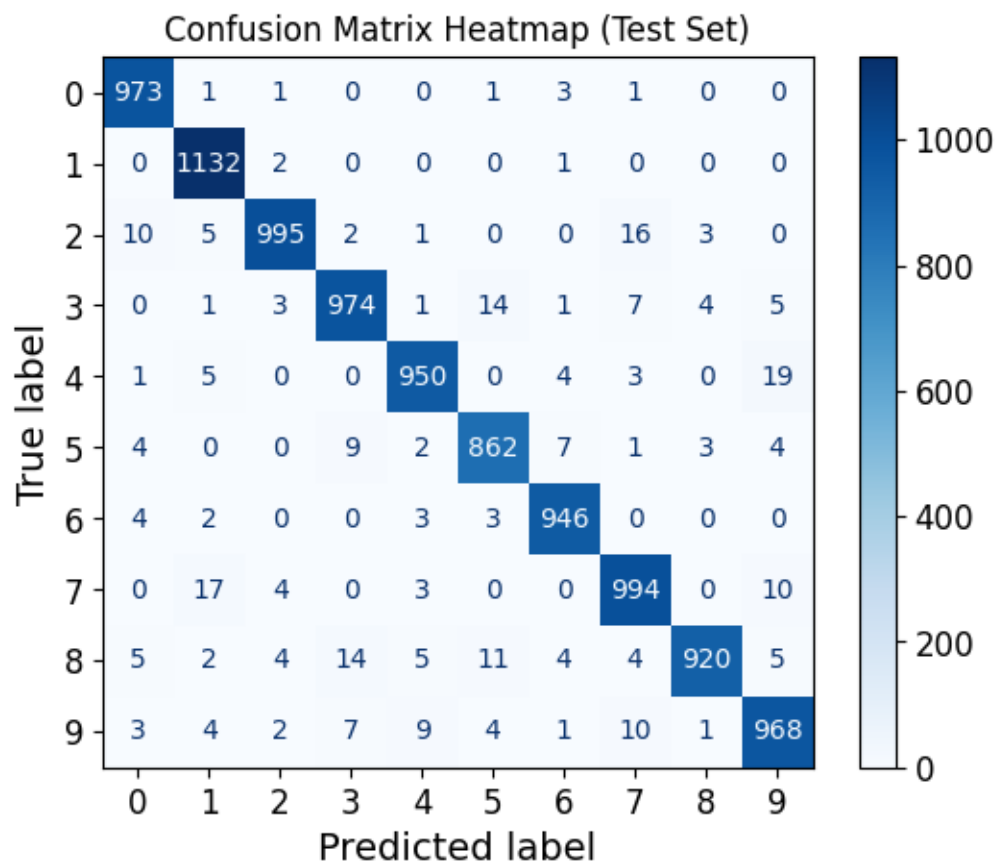
```
[61]: y_knn_expanded_pred = knn_clf.predict(X_test)
```

```
[62]: accuracy_score(y_test, y_knn_expanded_pred)
```

```
[62]: 0.9763
```

```
[68]: # prompt: create a confusionnmatric heatmap

cm = confusion_matrix(y_test, y_knn_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=knn_clf.
    ↪classes_)
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix Heatmap (Test Set)")
plt.show()
```



```
[69]: print("\n" + "*** Classification Report (Test Set) ***"*3)
print(classification_report(y_test, y_knn_pred))
```

```
*** Classification Report (Test Set) ***** Classification Report (Test Set)
***** Classification Report (Test Set) ***
           precision    recall  f1-score   support


```



0	0.97	0.99	0.98	980
1	0.97	1.00	0.98	1135
2	0.98	0.96	0.97	1032
3	0.97	0.96	0.97	1010
4	0.98	0.97	0.97	982
5	0.96	0.97	0.96	892
6	0.98	0.99	0.98	958
7	0.96	0.97	0.96	1028
8	0.99	0.94	0.97	974
9	0.96	0.96	0.96	1009
accuracy				0.97 10000
macro avg				0.97 10000
weighted avg				0.97 10000

```
[63]: ambiguous_digit = X_test[2589]
      knn_clf.predict_proba([ambiguous_digit])
```

```
[63]: array([[0.24579675, 0.          , 0.          , 0.          , 0.          ,
              0.          , 0.          , 0.          , 0.          , 0.75420325]])
```

```
[64]: plot_digit(ambiguous_digit)
```



```
[72]: !pip install gradio -q

import gradio as gr
import numpy as np
from PIL import Image

# Load the trained KNN classifier
knn_clf = KNeighborsClassifier(weights='distance', n_neighbors=4)
knn_clf.fit(X_train_expanded, y_train_expanded) # Using the expanded training
↳ set

def classify_digit(image):
    """Classifies a single grayscale image of a digit using the trained KNN model.
    ↳ """
    # Ensure the input is a NumPy array and is grayscale
    image = np.array(image).astype(np.uint8)
    # Ensure the image is 28x28
    if image.shape != (28, 28):
        # Resize the image if it's not 28x28 (optional, but good practice for
        ↳ consistency)
        image = Image.fromarray(image).resize((28, 28), Image.LANCZOS)
        image = np.array(image)

    # Flatten the image to a 1D array (784 features)
    image_flattened = image.reshape(1, -1)

    # Make a prediction
    prediction = knn_clf.predict(image_flattened)[0]

    return str(prediction)

# Create the Gradio interface
iface = gr.Interface(
    fn=classify_digit,
    inputs=gr.Image(width=28, height=28, image_mode='L'), # Grayscale image
    ↳ input
    outputs="text",
    title="MNIST Digit Classifier",
    description="Upload a grayscale image of a digit (0-9) and the model will
    ↳ classify it.",
    live=True # Enable live prediction as the user draws/uploads
)

# Launch the interface
iface.launch(debug=True)
```

It looks like you are running Gradio on a hosted Jupyter notebook. For the Gradio app to work, sharing must be enabled. Automatically setting `share=True` (you can turn this off by setting `share=False` in `launch()` explicitly).

Colab notebook detected. This cell will run indefinitely so that you can see errors and logs. To turn off, set debug=False in launch().

\* Running on public URL: <https://877d692049d404a873.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgrades, run `gradio deploy` from the terminal in the working directory to deploy to Hugging Face Spaces (<https://huggingface.co/spaces>)

<IPython.core.display.HTML object>

Created dataset file at: .gradio/flagged/dataset1.csv

Traceback (most recent call last):

File "/usr/local/lib/python3.11/dist-packages/gradio/queueing.py", line 625, in process\_events

response = await route\_utils.call\_process\_api(  
~~~~~

File "/usr/local/lib/python3.11/dist-packages/gradio/route\_utils.py", line 322, in call\_process\_api

output = await app.get\_blocks().process\_api(  
~~~~~

File "/usr/local/lib/python3.11/dist-packages/gradio/blocks.py", line 2191, in process\_api

result = await self.call\_function(  
~~~~~

File "/usr/local/lib/python3.11/dist-packages/gradio/blocks.py", line 1702, in call\_function

prediction = await anyio.to\_thread.run\_sync( # type: ignore  
~~~~~

File "/usr/local/lib/python3.11/dist-packages/anyio/to\_thread.py", line 56, in run\_sync

return await get\_async\_backend().run\_sync\_in\_worker\_thread(  
~~~~~

File "/usr/local/lib/python3.11/dist-packages/anyio/\_backends/\_asyncio.py", line 2470, in run\_sync\_in\_worker\_thread

return await future  
~~~~~

File "/usr/local/lib/python3.11/dist-packages/anyio/\_backends/\_asyncio.py", line 967, in run

result = context.run(func, \*args)  
~~~~~

File "/usr/local/lib/python3.11/dist-packages/gradio/utils.py", line 894, in wrapper

response = f(\*args, \*\*kwargs)  
~~~~~

```
File "/tmp/ipython-input-72-48946710.py", line 14, in classify_digit
    image = np.array(image).astype(np.uint8)
    ~~~~~
```

TypeError: int() argument must be a string, a bytes-like object or a real number, not 'NoneType'

Keyboard interruption in main thread... closing server.

Killing tunnel 127.0.0.1:7860 <> https://877d692049d404a873.gradio.live

[72]:

## 6 Exercise solutions

### 6.1 1. An MNIST Classifier With Over 97% Accuracy

```
[ ]: from sklearn.model_selection import GridSearchCV

param_grid = [{'weights': ["uniform", "distance"], 'n_neighbors': [3, 4, 5]}]

knn_clf = KNeighborsClassifier()
grid_search = GridSearchCV(knn_clf, param_grid, cv=5, verbose=3)
grid_search.fit(X_train, y_train)
```

Fitting 5 folds for each of 6 candidates, totalling 30 fits

[CV] n\_neighbors=3, weights=uniform ...

[Parallel(n\_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[CV] ... n\_neighbors=3, weights=uniform, score=0.972, total=168.0min

[CV] n\_neighbors=3, weights=uniform ...

[Parallel(n\_jobs=1)]: Done 1 out of 1 | elapsed: 168.0min remaining: 0.0s

[CV] ... n\_neighbors=3, weights=uniform, score=0.971, total=12.3min

[CV] n\_neighbors=3, weights=uniform ...

[Parallel(n\_jobs=1)]: Done 2 out of 2 | elapsed: 180.3min remaining: 0.0s

[CV] ... n\_neighbors=3, weights=uniform, score=0.969, total=11.9min

[CV] n\_neighbors=3, weights=uniform ...

[CV] ... n\_neighbors=3, weights=uniform, score=0.969, total=12.5min

[CV] n\_neighbors=3, weights=uniform ...

[CV] ... n\_neighbors=3, weights=uniform, score=0.970, total=12.7min

[CV] n\_neighbors=3, weights=distance ...

[CV] ... n\_neighbors=3, weights=distance, score=0.972, total=12.5min

[CV] n\_neighbors=3, weights=distance ...

[CV] ... n\_neighbors=3, weights=distance, score=0.972, total=12.8min

[CV] n\_neighbors=3, weights=distance ...

[CV] ... n\_neighbors=3, weights=distance, score=0.970, total=12.6min

[CV] n\_neighbors=3, weights=distance ...

[CV] ... n\_neighbors=3, weights=distance, score=0.970, total=12.9min

```

[CV] n_neighbors=3, weights=distance ...
[CV] ... n_neighbors=3, weights=distance, score=0.971, total=11.3min
[CV] n_neighbors=4, weights=uniform ...
[CV] ... n_neighbors=4, weights=uniform, score=0.969, total=11.0min
[CV] n_neighbors=4, weights=uniform ...
[CV] ... n_neighbors=4, weights=uniform, score=0.968, total=11.0min
[CV] n_neighbors=4, weights=uniform ...
[CV] ... n_neighbors=4, weights=uniform, score=0.968, total=11.0min
[CV] n_neighbors=4, weights=uniform ...
[CV] ... n_neighbors=4, weights=uniform, score=0.967, total=11.0min
[CV] n_neighbors=4, weights=uniform ...
[CV] ... n_neighbors=4, weights=uniform, score=0.970, total=11.0min
[CV] n_neighbors=4, weights=distance ...
[CV] ... n_neighbors=4, weights=distance, score=0.973, total=11.0min
[CV] n_neighbors=4, weights=distance ...
[CV] ... n_neighbors=4, weights=distance, score=0.972, total=11.0min
[CV] n_neighbors=4, weights=distance ...
[CV] ... n_neighbors=4, weights=distance, score=0.970, total=11.0min
[CV] n_neighbors=4, weights=distance ...
[CV] ... n_neighbors=4, weights=distance, score=0.971, total=11.0min
[CV] n_neighbors=4, weights=distance ...
[CV] ... n_neighbors=4, weights=distance, score=0.972, total=11.3min
[CV] n_neighbors=5, weights=uniform ...
[CV] ... n_neighbors=5, weights=uniform, score=0.970, total=10.9min
[CV] n_neighbors=5, weights=uniform ...
[CV] ... n_neighbors=5, weights=uniform, score=0.970, total=11.0min
[CV] n_neighbors=5, weights=uniform ...
[CV] ... n_neighbors=5, weights=uniform, score=0.969, total=11.0min
[CV] n_neighbors=5, weights=uniform ...
[CV] ... n_neighbors=5, weights=uniform, score=0.968, total=11.1min
[CV] n_neighbors=5, weights=uniform ...
[CV] ... n_neighbors=5, weights=uniform, score=0.969, total=11.0min
[CV] n_neighbors=5, weights=distance ...
[CV] ... n_neighbors=5, weights=distance, score=0.970, total=93.6min
[CV] n_neighbors=5, weights=distance ...
[CV] ... n_neighbors=5, weights=distance, score=0.971, total=11.0min
[CV] n_neighbors=5, weights=distance ...
[CV] ... n_neighbors=5, weights=distance, score=0.970, total=10.9min
[CV] n_neighbors=5, weights=distance ...
[CV] ... n_neighbors=5, weights=distance, score=0.969, total=11.2min
[CV] n_neighbors=5, weights=distance ...
[CV] ... n_neighbors=5, weights=distance, score=0.971, total=11.1min
[Parallel(n_jobs=1)]: Done 30 out of 30 | elapsed: 582.5min finished

```

```

[ ]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(),
                param_grid=[{'n_neighbors': [3, 4, 5],
                              'weights': ['uniform', 'distance']}],

```

```
verbose=3)
```

```
[ ]: grid_search.best_params_
```

```
[ ]: {'n_neighbors': 4, 'weights': 'distance'}
```

```
[ ]: grid_search.best_score_
```

```
[ ]: 0.9716166666666666
```

```
[ ]: from sklearn.metrics import accuracy_score
```

```
y_pred = grid_search.predict(X_test)
accuracy_score(y_test, y_pred)
```

```
[ ]: 0.9714
```


## 6.2 2. Data Augmentation

```
[ ]: from scipy.ndimage.interpolation import shift
```

```
[ ]: def shift_image(image, dx, dy):
    image = image.reshape((28, 28))
    shifted_image = shift(image, [dy, dx], cval=0, mode="constant")
    return shifted_image.reshape([-1])
```

```
[ ]: image = X_train[1000]
    shifted_image_down = shift_image(image, 0, 5)
    shifted_image_left = shift_image(image, -5, 0)

    plt.figure(figsize=(12,3))
    plt.subplot(131)
    plt.title("Original", fontsize=14)
    plt.imshow(image.reshape(28, 28), interpolation="nearest", cmap="Greys")
    plt.subplot(132)
    plt.title("Shifted down", fontsize=14)
    plt.imshow(shifted_image_down.reshape(28, 28), interpolation="nearest",
    ↪cmap="Greys")
    plt.subplot(133)
    plt.title("Shifted left", fontsize=14)
    plt.imshow(shifted_image_left.reshape(28, 28), interpolation="nearest",
    ↪cmap="Greys")
    plt.show()
```



```
[ ]: X_train_augmented = [image for image in X_train]
    y_train_augmented = [label for label in y_train]

    for dx, dy in ((1, 0), (-1, 0), (0, 1), (0, -1)):
        for image, label in zip(X_train, y_train):
            X_train_augmented.append(shift_image(image, dx, dy))
            y_train_augmented.append(label)

X_train_augmented = np.array(X_train_augmented)
y_train_augmented = np.array(y_train_augmented)
```

```
[ ]: shuffle_idx = np.random.permutation(len(X_train_augmented))
    X_train_augmented = X_train_augmented[shuffle_idx]
    y_train_augmented = y_train_augmented[shuffle_idx]
```

```
[ ]: knn_clf = KNeighborsClassifier(**grid_search.best_params_)
```

```
[ ]: knn_clf.fit(X_train_augmented, y_train_augmented)
```

```
[ ]: KNeighborsClassifier(n_neighbors=4, weights='distance')
```

**Warning:** the following cell may take close to an hour to run, depending on your hardware.

```
[ ]: y_pred = knn_clf.predict(X_test)
    accuracy_score(y_test, y_pred)
```

```
[ ]: 0.9763
```

By simply augmenting the data, we got a 0.5% accuracy boost. :)

### 6.3 3. Tackle the Titanic dataset

The goal is to predict whether or not a passenger survived based on attributes such as their age, sex, passenger class, where they embarked and so on.

Let's fetch the data and load it:

```
[ ]: import os
import urllib.request

TITANIC_PATH = os.path.join("datasets", "titanic")
DOWNLOAD_URL = "https://raw.githubusercontent.com/ageron/handson-ml2/master/
↳datasets/titanic/"

def fetch_titanic_data(url=DOWNLOAD_URL, path=TITANIC_PATH):
    if not os.path.isdir(path):
        os.makedirs(path)
    for filename in ("train.csv", "test.csv"):
        filepath = os.path.join(path, filename)
        if not os.path.isfile(filepath):
            print("Downloading", filename)
            urllib.request.urlretrieve(url + filename, filepath)

fetch_titanic_data()
```

```
[ ]: import pandas as pd

def load_titanic_data(filename, titanic_path=TITANIC_PATH):
    csv_path = os.path.join(titanic_path, filename)
    return pd.read_csv(csv_path)
```

```
[ ]: train_data = load_titanic_data("train.csv")
test_data = load_titanic_data("test.csv")
```

The data is already split into a training set and a test set. However, the test data does *not* contain the labels: your goal is to train the best model you can using the training data, then make your predictions on the test data and upload them to Kaggle to see your final score.

```
[ ]: train_data.head()
```

```
[ ]: 
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	



	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

Let's explicitly set the PassengerId column as the index column:

```
[ ]: train_data = train_data.set_index("PassengerId")
test_data = test_data.set_index("PassengerId")
```

Let's get more info to see how much data is missing:

```
[ ]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 891 entries, 1 to 891
Data columns (total 11 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Survived    891 non-null    int64
1   Pclass      891 non-null    int64
2   Name        891 non-null    object
3   Sex         891 non-null    object
4   Age         714 non-null    float64
5   SibSp       891 non-null    int64
6   Parch       891 non-null    int64
7   Ticket      891 non-null    object
8   Fare        891 non-null    float64
9   Cabin       204 non-null    object
10  Embarked    889 non-null    object
dtypes: float64(2), int64(4), object(5)
memory usage: 83.5+ KB
```

```
[ ]: train_data[train_data["Sex"]=="female"]["Age"].median()
```

```
[ ]: 27.0
```

Let's take a look at the numerical attributes:

```
[ ]: train_data.describe()
```

	Survived	Pclass	Age	SibSp	Parch	Fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699113	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526507	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.416700	0.000000	0.000000	0.000000

25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

Now let's build our preprocessing pipelines, starting with the pipeline for numerical attributes:

```
[ ]: from sklearn.pipeline import Pipeline
      from sklearn.impute import SimpleImputer
      from sklearn.preprocessing import StandardScaler

      num_pipeline = Pipeline([
          ("imputer", SimpleImputer(strategy="median")),
          ("scaler", StandardScaler())
      ])
```

Now we can build the pipeline for the categorical attributes:

```
[ ]: from sklearn.preprocessing import OneHotEncoder

[ ]: cat_pipeline = Pipeline([
      ("imputer", SimpleImputer(strategy="most_frequent")),
      ("cat_encoder", OneHotEncoder(sparse=False)),
  ])
```

Finally, let's join the numerical and categorical pipelines:

```
[ ]: from sklearn.compose import ColumnTransformer

      num_attribs = ["Age", "SibSp", "Parch", "Fare"]
      cat_attribs = ["Pclass", "Sex", "Embarked"]

      preprocess_pipeline = ColumnTransformer([
          ("num", num_pipeline, num_attribs),
          ("cat", cat_pipeline, cat_attribs),
      ])
```

Cool! Now we have a nice preprocessing pipeline that takes the raw data and outputs numerical input features that we can feed to any Machine Learning model we want.

```
[ ]: X_train = preprocess_pipeline.fit_transform(
      train_data[num_attribs + cat_attribs])
      X_train
```

```
[ ]: array([[ -0.56573582,  0.43279337, -0.47367361, ...,  0.          ,
              0.          ,  1.          ],
            [ 0.6638609 ,  0.43279337, -0.47367361, ...,  1.          ,
              0.          ,  0.          ]],
```

```

[-0.25833664, -0.4745452 , -0.47367361, ..., 0.          ,
 0.          , 1.          ],
...,
[-0.10463705, 0.43279337, 2.00893337, ..., 0.          ,
 0.          , 1.          ],
[-0.25833664, -0.4745452 , -0.47367361, ..., 1.          ,
 0.          , 0.          ],
[ 0.20276213, -0.4745452 , -0.47367361, ..., 0.          ,
 1.          , 0.          ]])

```

```
[ ]: y_train = train_data["Survived"]
```

```
[ ]: from sklearn.ensemble import RandomForestClassifier

forest_clf = RandomForestClassifier(n_estimators=100, random_state=42)
forest_clf.fit(X_train, y_train)
```

```
[ ]: RandomForestClassifier(random_state=42)
```

```
[ ]: X_test = preprocess_pipeline.transform(test_data[num_attribs + cat_attribs])
y_pred = forest_clf.predict(X_test)
```

```
[ ]: from sklearn.model_selection import cross_val_score

forest_scores = cross_val_score(forest_clf, X_train, y_train, cv=10)
forest_scores.mean()
```

```
[ ]: 0.8137578027465668
```


```
[ ]: from sklearn.svm import SVC

svm_clf = SVC(gamma="auto")
svm_scores = cross_val_score(svm_clf, X_train, y_train, cv=10)
svm_scores.mean()
```

```
[ ]: 0.8249313358302123
```

```
[ ]: import matplotlib.pyplot as plt

plt.figure(figsize=(8, 4))
plt.plot([1]*10, svm_scores, ".")
plt.plot([2]*10, forest_scores, ".")
plt.boxplot([svm_scores, forest_scores], labels=("SVM", "Random Forest"))
plt.ylabel("Accuracy", fontsize=14)
plt.show()
```



```
[ ]: train_data["AgeBucket"] = train_data["Age"] // 15 * 15
train_data[["AgeBucket", "Survived"]].groupby(['AgeBucket']).mean()
```

```
[ ]:      Survived
AgeBucket
0.0      0.576923
15.0     0.362745
30.0     0.423256
45.0     0.404494
60.0     0.240000
75.0     1.000000
```

```
[ ]: train_data["RelativesOnboard"] = train_data["SibSp"] + train_data["Parch"]
train_data[["RelativesOnboard", "Survived"]].groupby(['RelativesOnboard']).
    ↪mean()
```

```
[ ]:      Survived
RelativesOnboard
0      0.303538
1      0.552795
2      0.578431
3      0.724138
4      0.200000
5      0.136364
6      0.333333
7      0.000000
10     0.000000
```

## 6.4 4. Spam classifier

First, let's fetch the data:

```
[ ]: import os
import tarfile
import urllib.request

DOWNLOAD_ROOT = "http://spamassassin.apache.org/old/publiccorpus/"
HAM_URL = DOWNLOAD_ROOT + "20030228_easy_ham.tar.bz2"
SPAM_URL = DOWNLOAD_ROOT + "20030228_spam.tar.bz2"
SPAM_PATH = os.path.join("datasets", "spam")

def fetch_spam_data(ham_url=HAM_URL, spam_url=SPAM_URL, spam_path=SPAM_PATH):
    if not os.path.isdir(spam_path):
        os.makedirs(spam_path)
    for filename, url in (("ham.tar.bz2", ham_url), ("spam.tar.bz2", spam_url)):
        path = os.path.join(spam_path, filename)
        if not os.path.isfile(path):
            urllib.request.urlretrieve(url, path)
            tar_bz2_file = tarfile.open(path)
            tar_bz2_file.extractall(path=spam_path)
            tar_bz2_file.close()
```

```
[ ]: fetch_spam_data()
```

Next, let's load all the emails:

```
[ ]: HAM_DIR = os.path.join(SPAM_PATH, "easy_ham")
SPAM_DIR = os.path.join(SPAM_PATH, "spam")
ham_filenames = [name for name in sorted(os.listdir(HAM_DIR)) if len(name) > 20]
spam_filenames = [name for name in sorted(os.listdir(SPAM_DIR)) if len(name) > 20]
```

```
[ ]: len(ham_filenames)
```

```
[ ]: 2500
```

```
[ ]: len(spam_filenames)
```

```
[ ]: 500
```

We can use Python's `email` module to parse these emails (this handles headers, encoding, and so on):

```
[ ]: import email
import email.policy

def load_email(is_spam, filename, spam_path=SPAM_PATH):
```

```

directory = "spam" if is_spam else "easy_ham"
with open(os.path.join(spam_path, directory, filename), "rb") as f:
    return email.parser.BytesParser(policy=email.policy.default).parse(f)

```

```

[ ]: ham_emails = [load_email(is_spam=False, filename=name) for name in
    ↳ ham_filenames]
spam_emails = [load_email(is_spam=True, filename=name) for name in
    ↳ spam_filenames]

```

Let's look at one example of ham and one example of spam, to get a feel of what the data looks like:

```

[ ]: print(ham_emails[1].get_content().strip())

```

Martin A posted:

Tassos Papadopoulos, the Greek sculptor behind the plan, judged that the limestone of Mount Kerdyllo, 70 miles east of Salonika and not far from the Mount Athos monastic community, was ideal for the patriotic sculpture.

As well as Alexander's granite features, 240 ft high and 170 ft wide, a museum, a restored amphitheatre and car park for admiring crowds are planned

-----

So is this mountain limestone or granite?  
If it's limestone, it'll weather pretty fast.

----- Yahoo! Groups Sponsor ----->  
4 DVDs Free +s&p Join Now  
<http://us.click.yahoo.com/pt6YBB/NXiEAA/mG3HAA/7gSolB/TM>  
----->

To unsubscribe from this group, send an email to:  
forteana-unsubscribe@egroups.com

Your use of Yahoo! Groups is subject to <http://docs.yahoo.com/info/terms/>

```

[ ]: print(spam_emails[6].get_content().strip())

```

Help wanted. We are a 14 year old fortune 500 company, that is growing at a tremendous rate. We are looking for individuals who want to work from home.

This is an opportunity to make an excellent income. No experience is required. We will train you.

So if you are looking to be employed from home with a career that has

vast opportunities, then go:

<http://www.basetel.com/wealthnow>

We are looking for energetic and self motivated people. If that is you than click on the link and fill out the form, and one of our employment specialist will contact you.

To be removed from our link simple go to:

<http://www.basetel.com/remove.html>

4139vOLW7-758DoDY1425FRhM1-764SMFc8513fCsL140

Some emails are actually multipart, with images and attachments (which can have their own attachments). Let's look at the various types of structures we have:

```
[ ]: def get_email_structure(email):
    if isinstance(email, str):
        return email
    payload = email.get_payload()
    if isinstance(payload, list):
        return "multipart({})".format(", ".join([
            get_email_structure(sub_email)
            for sub_email in payload
        ]))
    else:
        return email.get_content_type()
```

```
[ ]: from collections import Counter

def structures_counter(emails):
    structures = Counter()
    for email in emails:
        structure = get_email_structure(email)
        structures[structure] += 1
    return structures
```

```
[ ]: structures_counter(ham_emails).most_common()
```

```
[ ]: [('text/plain', 2408),
      ('multipart(text/plain, application/pgp-signature)', 66),
      ('multipart(text/plain, text/html)', 8),
      ('multipart(text/plain, text/plain)', 4),
      ('multipart(text/plain)', 3),
      ('multipart(text/plain, application/octet-stream)', 2),
      ('multipart(text/plain, text/enriched)', 1),
```

```
(('multipart(text/plain, application/ms-tnef, text/plain)', 1),
 ('multipart(multipart(text/plain, text/plain, text/plain), application/pgp-
signature)',
 1),
 ('multipart(text/plain, video/mng)', 1),
 ('multipart(text/plain, multipart(text/plain))', 1),
 ('multipart(text/plain, application/x-pkcs7-signature)', 1),
 ('multipart(text/plain, multipart(text/plain, text/plain),
text/rfc822-headers)',
 1),
 ('multipart(text/plain, multipart(text/plain, text/plain),
multipart(multipart(text/plain, application/x-pkcs7-signature)))',
 1),
 ('multipart(text/plain, application/x-java-applet)', 1)]
```

```
[ ]: structures_counter(spam_emails).most_common()
```

```
[ ]: [('text/plain', 218),
 ('text/html', 183),
 ('multipart(text/plain, text/html)', 45),
 ('multipart(text/html)', 20),
 ('multipart(text/plain)', 19),
 ('multipart(multipart(text/html))', 5),
 ('multipart(text/plain, image/jpeg)', 3),
 ('multipart(text/html, application/octet-stream)', 2),
 ('multipart(text/plain, application/octet-stream)', 1),
 ('multipart(text/html, text/plain)', 1),
 ('multipart(multipart(text/html), application/octet-stream, image/jpeg)', 1),
 ('multipart(multipart(text/plain, text/html), image/gif)', 1),
 ('multipart/alternative', 1)]
```

It seems that the ham emails are more often plain text, while spam has quite a lot of HTML. Moreover, quite a few ham emails are signed using PGP, while no spam is. In short, it seems that the email structure is useful information to have.

Now let's take a look at the email headers:

```
[ ]: for header, value in spam_emails[0].items():
      print(header, ":", value)
```

```
Return-Path : <12a1mailbot1@web.de>
Delivered-To : zzzz@localhost.spamassassin.taint.org
Received : from localhost (localhost [127.0.0.1])          by
phobos.labs.spamassassin.taint.org (Postfix) with ESMTP id 136B943C32          for
<zzzz@localhost>; Thu, 22 Aug 2002 08:17:21 -0400 (EDT)
Received : from mail.webnote.net [193.120.211.219]          by localhost with POP3
(fetchmail-5.9.0)          for zzzz@localhost (single-drop); Thu, 22 Aug 2002
13:17:21 +0100 (IST)
```



```
Received : from dd_it7 ([210.97.77.167])          by webnote.net (8.9.3/8.9.3)
with ESMTP id NAA04623      for <zzzz@spamassassin.taint.org>; Thu, 22 Aug 2002
13:09:41 +0100
From : 12a1mailbot1@web.de
Received : from r-smtp.korea.com - 203.122.2.197 by dd_it7  with Microsoft
SMTPSVC(5.5.1775.675.6);      Sat, 24 Aug 2002 09:42:10 +0900
To : dcek1a1@netsgo.com
Subject : Life Insurance - Why Pay More?
Date : Wed, 21 Aug 2002 20:31:57 -1600
MIME-Version : 1.0
Message-ID : <0103c1042001882DD_IT7@dd_it7>
Content-Type : text/html; charset="iso-8859-1"
Content-Transfer-Encoding : quoted-printable
```

There's probably a lot of useful information in there, such as the sender's email address (12a1mailbot1@web.de looks fishy), but we will just focus on the Subject header:

```
[ ]: spam_emails[0]["Subject"]
```

```
[ ]: 'Life Insurance - Why Pay More?'
```

Okay, before we learn too much about the data, let's not forget to split it into a training set and a test set:

```
[ ]: import numpy as np
from sklearn.model_selection import train_test_split

X = np.array(ham_emails + spam_emails, dtype=object)
y = np.array([0] * len(ham_emails) + [1] * len(spam_emails))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[ ]: import re
from html import unescape

def html_to_plain_text(html):
    text = re.sub('<head.*?>.*?</head>', '', html, flags=re.M | re.S | re.I)
    text = re.sub('<a\s.*?>', ' HYPERLINK ', text, flags=re.M | re.S | re.I)
    text = re.sub('<.*?>', '', text, flags=re.M | re.S)
    text = re.sub(r'(\s*\n)+', '\n', text, flags=re.M | re.S)
    return unescape(text)
```

Let's see if it works. This is HTML spam:

```
[ ]: html_spam_emails = [email for email in X_train[y_train==1]
                        if get_email_structure(email) == "text/html"]
sample_html_spam = html_spam_emails[7]
print(sample_html_spam.get_content().strip()[:1000], "...")
```

```
[ ]: print(html_to_plain_text(sample_html_spam.get_content())[:1000], "...")
```

```
[ ]: def email_to_text(email):  
    html = None  
    for part in email.walk():  
        ctype = part.get_content_type()  
        if not ctype in ("text/plain", "text/html"):  
            continue
```

```

try:
    content = part.get_content()
except: # in case of encoding issues
    content = str(part.get_payload())
if ctype == "text/plain":
    return content
else:
    html = content
if html:
    return html_to_plain_text(html)

```

```
[ ]: print(email_to_text(sample_html_spam)[:100], "...")
```

```

OTC
Newsletter
Discover Tomorrow's Winners
For Immediate Release
Cal-Bay (Stock Symbol: CBYI)
Wat ...

```

```

[ ]: try:
    import nltk

    stemmer = nltk.PorterStemmer()
    for word in ("Computations", "Computation", "Computing", "Computed", "Compute", "Compulsive"):
        print(word, "=>", stemmer.stem(word))
except ImportError:
    print("Error: stemming requires the NLTK module.")
    stemmer = None

```

```

Computations => comput
Computation => comput
Computing => comput
Computed => comput
Compute => comput
Compulsive => compuls

```

We will also need a way to replace URLs with the word “URL”. For this, we could use hard core [regular expressions](#) but we will just use the [urlextract](#) library. You can install it with the following command (don’t forget to activate your virtualenv first; if you don’t have one, you will likely need administrator rights, or use the `--user` option):

```
$ pip3 install urlextract
```

```

[ ]: # if running this notebook on Colab or Kaggle, we just pip install urlextract
if IS_COLAB or IS_KAGGLE:

```

```
%pip install -q -U urlextract
```

```
[ ]: try:
    import urlextract # may require an Internet connection to download root
    ↪ domain names

    url_extractor = urlextract.URLExtract()
    print(url_extractor.find_urls("Will it detect github.com and https://youtu.
    ↪ be/7Pq-S557XQU?t=3m32s"))
except ImportError:
    print("Error: replacing URLs requires the urlextract module.")
    url_extractor = None
```

```
['github.com', 'https://youtu.be/7Pq-S557XQU?t=3m32s']
```

```
[ ]: from sklearn.base import BaseEstimator, TransformerMixin

class EmailToWordCounterTransformer(BaseEstimator, TransformerMixin):
    def __init__(self, strip_headers=True, lower_case=True,
    ↪ remove_punctuation=True,
        replace_urls=True, replace_numbers=True, stemming=True):
        self.strip_headers = strip_headers
        self.lower_case = lower_case
        self.remove_punctuation = remove_punctuation
        self.replace_urls = replace_urls
        self.replace_numbers = replace_numbers
        self.stemming = stemming
    def fit(self, X, y=None):
        return self
    def transform(self, X, y=None):
        X_transformed = []
        for email in X:
            text = email_to_text(email) or ""
            if self.lower_case:
                text = text.lower()
            if self.replace_urls and url_extractor is not None:
                urls = list(set(url_extractor.find_urls(text)))
                urls.sort(key=lambda url: len(url), reverse=True)
                for url in urls:
                    text = text.replace(url, " URL ")
            if self.replace_numbers:
                text = re.sub(r'\d+(?:\.\d*)?(?:[eE][+-]?\d+)?', 'NUMBER', text)
            if self.remove_punctuation:
                text = re.sub(r'\W+', ' ', text, flags=re.M)
            word_counts = Counter(text.split())
            if self.stemming and stemmer is not None:
                stemmed_word_counts = Counter()
```

```

        for word, count in word_counts.items():
            stemmed_word = stemmer.stem(word)
            stemmed_word_counts[stemmed_word] += count
        word_counts = stemmed_word_counts
        X_transformed.append(word_counts)
    return np.array(X_transformed)

```

Let's try this transformer on a few emails:

```

[ ]: X_few = X_train[:3]
X_few_wordcounts = EmailToWordCounterTransformer().fit_transform(X_few)
X_few_wordcounts

[ ]: array([Counter({'chuck': 1, 'murcko': 1, 'wrote': 1, 'stuff': 1, 'yawn': 1, 'r':
1}),
          Counter({'the': 11, 'of': 9, 'and': 8, 'all': 3, 'christian': 3, 'to': 3,
'by': 3, 'jefferson': 2, 'i': 2, 'have': 2, 'superstit': 2, 'one': 2, 'on': 2,
'been': 2, 'ha': 2, 'half': 2, 'rogueri': 2, 'teach': 2, 'jesu': 2, 'some': 1,
'interest': 1, 'quot': 1, 'url': 1, 'thoma': 1, 'examin': 1, 'known': 1, 'word':
1, 'do': 1, 'not': 1, 'find': 1, 'in': 1, 'our': 1, 'particular': 1, 'redeem':
1, 'featur': 1, 'they': 1, 'are': 1, 'alike': 1, 'found': 1, 'fabl': 1,
'mytholog': 1, 'million': 1, 'innoc': 1, 'men': 1, 'women': 1, 'children': 1,
'sinc': 1, 'introduct': 1, 'burnt': 1, 'tortur': 1, 'fine': 1, 'imprison': 1,
'what': 1, 'effect': 1, 'thi': 1, 'coercion': 1, 'make': 1, 'world': 1, 'fool':
1, 'other': 1, 'hypocrit': 1, 'support': 1, 'error': 1, 'over': 1, 'earth': 1,
'six': 1, 'histor': 1, 'american': 1, 'john': 1, 'e': 1, 'remsburg': 1,
'letter': 1, 'william': 1, 'short': 1, 'again': 1, 'becom': 1, 'most': 1,
'pervert': 1, 'system': 1, 'that': 1, 'ever': 1, 'shone': 1, 'man': 1, 'absurd':
1, 'untruth': 1, 'were': 1, 'perpetr': 1, 'upon': 1, 'a': 1, 'larg': 1, 'band':
1, 'dupe': 1, 'import': 1, 'led': 1, 'paul': 1, 'first': 1, 'great': 1,
'corrupt': 1}),
          Counter({'url': 4, 's': 3, 'group': 3, 'to': 3, 'in': 2, 'fortean': 2,
'martin': 2, 'an': 2, 'and': 2, 'we': 2, 'is': 2, 'yahoo': 2, 'unsubscribe': 2,
'y': 1, 'adamson': 1, 'wrote': 1, 'for': 1, 'altern': 1, 'rather': 1, 'more': 1,
'factual': 1, 'base': 1, 'rundown': 1, 'on': 1, 'hamza': 1, 'career': 1,
'includ': 1, 'hi': 1, 'belief': 1, 'that': 1, 'all': 1, 'non': 1, 'muslim': 1,
'yemen': 1, 'should': 1, 'be': 1, 'murder': 1, 'outright': 1, 'know': 1, 'how':
1, 'unbias': 1, 'memri': 1, 'don': 1, 't': 1, 'html': 1, 'rob': 1, 'sponsor': 1,
'number': 1, 'dvd': 1, 'free': 1, 'p': 1, 'join': 1, 'now': 1, 'from': 1, 'thi':
1, 'send': 1, 'email': 1, 'egroup': 1, 'com': 1, 'your': 1, 'use': 1, 'of': 1,
'subject': 1})],
          dtype=object)

```

```

[ ]: from scipy.sparse import csr_matrix

class WordCounterToVectorTransformer(BaseEstimator, TransformerMixin):
    def __init__(self, vocabulary_size=1000):

```

```

        self.vocabulary_size = vocabulary_size
    def fit(self, X, y=None):
        total_count = Counter()
        for word_count in X:
            for word, count in word_count.items():
                total_count[word] += min(count, 10)
        most_common = total_count.most_common()[:self.vocabulary_size]
        self.vocabulary_ = {word: index + 1 for index, (word, count) in
↪ enumerate(most_common)}
        return self
    def transform(self, X, y=None):
        rows = []
        cols = []
        data = []
        for row, word_count in enumerate(X):
            for word, count in word_count.items():
                rows.append(row)
                cols.append(self.vocabulary_.get(word, 0))
                data.append(count)
        return csr_matrix((data, (rows, cols)), shape=(len(X), self.
↪ vocabulary_size + 1))

```

```

[ ]: vocab_transformer = WordCounterToVectorTransformer(vocabulary_size=10)
X_few_vectors = vocab_transformer.fit_transform(X_few_wordcounts)
X_few_vectors

```

```

[ ]: <3x11 sparse matrix of type '<class 'numpy.longlong'>'
      with 20 stored elements in Compressed Sparse Row format>

```

```

[ ]: X_few_vectors.toarray()

```

```

[ ]: array([[ 6,  0,  0,  0,  0,  0,  0,  0,  0,  0,  0],
           [99, 11,  9,  8,  3,  1,  3,  1,  3,  2,  3],
           [67,  0,  1,  2,  3,  4,  1,  2,  0,  1,  0]], dtype=int64)

```

```

[ ]: vocab_transformer.vocabulary_

```

```

[ ]: {'the': 1,
      'of': 2,
      'and': 3,
      'to': 4,
      'url': 5,
      'all': 6,
      'in': 7,
      'christian': 8,
      'on': 9,
      'by': 10}

```

We are now ready to train our first spam classifier! Let's transform the whole dataset:

```
[ ]: from sklearn.pipeline import Pipeline

preprocess_pipeline = Pipeline([
    ("email_to_wordcount", EmailToWordCounterTransformer()),
    ("wordcount_to_vector", WordCounterToVectorTransformer()),
])

X_train_transformed = preprocess_pipeline.fit_transform(X_train)
```

**Note:** to be future-proof, we set solver="lbfgs" since this will be the default value in Scikit-Learn 0.22.

```
[ ]: from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import cross_val_score

log_clf = LogisticRegression(solver="lbfgs", max_iter=1000, random_state=42)
score = cross_val_score(log_clf, X_train_transformed, y_train, cv=3, verbose=3)
score.mean()
```

[Parallel(n\_jobs=1)]: Using backend SequentialBackend with 1 concurrent workers.

[Parallel(n\_jobs=1)]: Done 1 out of 1 | elapsed: 0.1s remaining: 0.0s

[CV] ...

[CV] ... , score=0.981, total= 0.1s

[CV] ...

[CV] ... , score=0.985, total= 0.2s

[CV] ...

[CV] ... , score=0.991, total= 0.2s

[Parallel(n\_jobs=1)]: Done 2 out of 2 | elapsed: 0.3s remaining: 0.0s

[Parallel(n\_jobs=1)]: Done 3 out of 3 | elapsed: 0.5s finished

```
[ ]: 0.9858333333333333
```

```
[ ]: from sklearn.metrics import precision_score, recall_score

X_test_transformed = preprocess_pipeline.transform(X_test)

log_clf = LogisticRegression(solver="lbfgs", max_iter=1000, random_state=42)
log_clf.fit(X_train_transformed, y_train)

y_pred = log_clf.predict(X_test_transformed)

print("Precision: {:.2f}%".format(100 * precision_score(y_test, y_pred)))
print("Recall: {:.2f}%".format(100 * recall_score(y_test, y_pred)))
```

Precision: 95.88%

Recall: 97.89%

[ ]: